



UvA-DARE (Digital Academic Repository)

Fireworks Algorithm versus Plant Propagation Algorithm

Vrieling, W.; van den Berg, D.

DOI

[10.5220/0008169401010112](https://doi.org/10.5220/0008169401010112)

Publication date

2019

Document Version

Final published version

Published in

IJCCI 2019

License

CC BY-NC-ND

[Link to publication](#)

Citation for published version (APA):

Vrieling, W., & van den Berg, D. (2019). Fireworks Algorithm versus Plant Propagation Algorithm. In J. J. Merelo, J. Garibaldi, A. Linares Barranco, K. Madani, & K. Warwick (Eds.), *IJCCI 2019: proceedings of the 11th International Joint Conference on Computational Intelligence : Vienna, Austria, September 17-19, 2019* (pp. 101-112). SciTePress Science and Technology Publications. <https://doi.org/10.5220/0008169401010112>

General rights

It is not permitted to download or to forward/distribute the text or part of it without the consent of the author(s) and/or copyright holder(s), other than for strictly personal, individual use, unless the work is under an open content license (like Creative Commons).

Disclaimer/Complaints regulations

If you believe that digital publication of certain material infringes any of your rights or (privacy) interests, please let the Library know, stating your reasons. In case of a legitimate complaint, the Library will make the material inaccessible and/or remove it from the website. Please Ask the Library: <https://uba.uva.nl/en/contact>, or a letter to: Library of the University of Amsterdam, Secretariat, Singel 425, 1012 WP Amsterdam, The Netherlands. You will be contacted as soon as possible.

Fireworks Algorithm versus Plant Propagation Algorithm

Wouter Vrieling¹ and Daan van den Berg²

¹*Computational Science Lab., Universiteit van Amsterdam, Science Park 904, Amsterdam, The Netherlands*

²*Institute for Informatics, Universiteit van Amsterdam, Science Park 904, Amsterdam, The Netherlands*

Keywords: Plant Propagation Algorithm, Fireworks Algorithm, Evolutionary Algorithm, Optimization Algorithm, Comparison, Metaheuristics.

Abstract: In recent years, the field of Evolutionary Algorithms has seen a tremendous increase in novel methods. While these algorithmic innovations often show excellent results on relatively limited domains, they are less often rigorously cross-tested or compared to other state-of-the-art developments. Two of these methods, quite similar in their appearance, are the Fireworks Algorithm and Plant Propagation Algorithm.

This study compares the similarities and differences between these two algorithms, from both quantitative and qualitative perspectives, by comparing them on a set of carefully chosen benchmark functions. The Fireworks Algorithm outperforms the Plant Propagation Algorithm on the majority of these, but when the functions are shifted slightly, Plant Propagation gives better results. Reasons behind these surprising differences are presented, and comparison methods for evolutionary algorithms are discussed in a wider context. All source code, graphs, test functions, and algorithmic implementations have been made publicly available for reference and further reuse.

1 INTRODUCTION

The increasingly popular field of population-based metaheuristics aims to create assumption-free and flexible algorithms. These Population-Based Algorithms (PBA) are transferable to different problem domains, easily implemented, almost effortlessly combined with exact methods, and capable of finding multiple solutions in a single run. Essential in the development of optimization algorithms is to balance *exploration* and *exploitation*. Exploration refers to the ability to sample the entire problem domain, ensuring that multiple options are considered. Exploitation refers to the ability to converge to a local optimum once in its basin, which is reflected directly in the quality of the solution found by the algorithm. It is important to correctly balance these properties to prevent getting stuck in local optima, or from never finding a neighbouring solution that might be better (Alba and Dorronsoro, 2005; Audibert et al., 2009; Ishii et al., 2002).

The Fireworks Algorithm (FWA) (Tan and Zhu, 2010) and the Plant Propagation Algorithm (PPA) (Salhi and Fraga, 2011) share a lot of algorithmic similarities and experimental results. Both are population-based algorithms that rely on normalizing fitness values of the individuals to (0,1). In both algorithms, relatively fit individuals generate many offspring with

small mutations, whereas relatively unfit individuals generate few offspring with large mutations. As such – even though the procedural details differ – both algorithms attempt to balance exploration and exploitation through relating the number of offspring and the degree of mutation to the normalized fitness of its parents.

Both FWA and PPA normalize fitness values; the best individual within the current population has a fitness of 1, whereas the worst individual within the current population has a fitness of 0. This makes the algorithms insensitive to linear changes or translations of the range of the objective function, as the relative difference between individuals does not depend on this range. The number of algorithmic parameters is reduced and less domain-specific assumptions have to be made, which in turn reduces the knowledge required about the algorithm and the effort required to find optimal algorithmic parameters.

The principle of reducing the number of algorithmic parameters can also be identified in the procedure that is responsible for the generation of offspring. In both algorithms, offspring are generated from a single parent. This eliminates the need for a crossover method, which – depending on the method used or the problem the algorithm is applied to – could be a source of constraint violation, requiring subsequent

repair methods and adding unnecessary complexity. However, this could be at the expense of combining good properties from multiple individuals into one and therefore could reduce exploitative properties of the algorithms (Paauw and van den Berg, 2019). Nonetheless, these principles create algorithms that are more widely applicable, as is so aptly described by Kenneth Sörensen (Sörensen, 2015). Many “novel” methods seem unnecessarily complex, are rarely extensively tested, or may introduce undesirable biases (Sörensen, 2015; Weyland, 2015).

Today, a quantitative comparison and qualitative treatise on the differences and similarities between the Fireworks Algorithm and the Plant Propagation Algorithm is presented. Strengths and weaknesses of both algorithms will be uncovered and performance analyses will be made to understand exactly how these algorithms differ and to what extent methods can be used to reduce the number of meta-parameters while increasing performance.

1.1 Related Work

Population-based algorithms are proving to be very successful in complex optimization problems and can occasionally overcome the drawbacks of traditional mathematical methods. The primary feature of PBA's is that, unlike heuristic methods such as Simulated Annealing (SA) (Kirkpatrick et al., 1983) and Tabu Search (TS) (Glover, 1989, 1990), evaluation of many candidate solutions is done simultaneously. The concurrent evaluation of many solutions enables the exploration of the search-space utilizing the collective intelligence that arises from these solutions, potentially avoiding local minima.

A classic success story of a Genetic Algorithm (GA) is that of reducing vibrations (and weight) of a space satellite boom structure (Moshrefi-Torbati et al., 2003; Keane, 1996). The GA significantly improved an existing design to 30 dB of energy transmission isolation between both ends. It has even actually been built. However, the evaluation procedure for this multi-objective optimization problem comes with great computational costs (El-Beltagy and Keane, 2001).

One older example of a variant of a PBA that is still often referred to is Differential Evolution (DE) (Storn and Price, 1997). The algorithm was developed with four main requirements in mind: the ability to handle different types of objective functions, parallelizability, ease of use (through a small number of control variables that are easy to choose), and consistent convergence properties. The main principle of DE is to add the weighted difference between two objective function's parameter vectors to a randomly

chosen third vector. This algorithm is very simple to implement and at the time of publishing (1997) outperformed many of the then prevalent strategies.

Another nature-inspired heuristic optimization algorithm that is often used in comparisons of optimization algorithms is the Harmony Search Algorithm (HSA) (Geem et al., 2001). It attempts to mimic the improvisation of musicians by evolving what it calls a set of “harmonies” that in turn consist of “pitches”¹. The main difference between HSA and GA is that HSA considers more than two parents when generating offspring and that it requires no careful choice of initial values. HSA shows good results for both discrete and continuous optimization problems, but is also criticized for not offering any novelty, apart from using a different terminology (Weyland, 2015).

The algorithms discussed in this study (FWA (Tan and Zhu, 2010) and PPA (Salhi and Fraga, 2011)) were published within such a narrow timeframe that the submissions might have easily crossed one another. However, FWA seems to be much more popular in the community than PPA. At the time of writing, more than 50 papers with “Fireworks Algorithm” in the title are available on Google Scholar, at least 10 of which have more than 50 citations. PPA has around ten papers with “Plant Propagation Algorithm” in their titles on Google Scholar, which have also been cited less often. A rough estimate would be that FWA is around five times bigger, in numbers of papers, in numbers of citations and possibly in the number of involved researchers.

The paper by Imran and Kowsalya (2014) is an interesting example of a real-world optimization problem that is solved by applying FWA. It shows a method using FWA to solve the problem of power system network reconfiguration in such a way that power loss and voltage profile is minimized. In this paper, FWA is compared with other classical methods from literature (GA, Refined GA, Improved TS, and HSA) and it is shown that FWA outperforms these methods in general quality of solutions.

Likewise, a recent example of PPA being applied to several real-world constrained optimization problems can be found in Sulaiman et al. (2014b). Herein PPA is found to be as good, or in some cases superior to, state-of-the-art optimization algorithms. More recently, a discrete version of PPA has been shown to perform admirably on the Traveling Salesman Problem (TSP), where it outperformed GA and SA both in solution quality and in runtime (Selamoğlu and Salhi, 2016). Interestingly, like Sörensen's paper, this paper also

¹Harmonies can be interpreted as individuals, while pitches can be interpreted as distinct discrete values of one parameter inside its domain.

poses the idea of comparing algorithms or heuristics on a number of other criteria than just performance; “[the criterion consisting of the number of arbitrarily set parameters required by a given heuristic algorithm is interesting because it does not depend on the programming skills of the researcher, the quality of code, the language of coding, the compiler or the processor. It is intrinsic to the algorithm itself.]” (Selamoğlu and Salhi, 2016).

Finally, in Cheraitia et al. (2017), a hybrid version of PPA and Local Search (LS) is used to solve the uncapacitated exam scheduling problem (UESP). Similarly, Geleijn et al. (2019) presents a non-hybrid version of PPA to solve the UESP. The UESP is a well-known computationally intractable combinatorial optimization problem that aims to schedule exams to periods while avoiding conflicts and spreading exams as evenly as possible. The papers show good results and demonstrate that PPA can easily be adapted to different types of problems.

1.2 Adaptations in Later Papers

EFWA (Enhanced FWA), as presented in Zheng et al. (2013), proposes five improvements over FWA: a new minimal explosion amplitude, a new method of generating explosion sparks, a new mapping strategy for sparks which are out of bounds, a new operator for generating Gaussian sparks and a new method for selection of individuals. EFWA not only outperforms FWA in convergence, but it also reduces runtime significantly (Zheng et al., 2013).

An adaptive approach (AFWA) is proposed for EFWA, wherein the formulas that determine the magnitude of mutation with which a new individual is generated are replaced with adaptive versions (Li et al., 2014). In this study, the distance at which new individuals are generated depends on information gained during the last iteration. The selected distance is proportional to the distance between the best new offspring and the closest offspring that is worse than the parent, as this is the area that will most likely hold better results. This results in performance improvement over EFWA, while it does not significantly increase computational costs.

Similarly, a dynamic approach (DynFWA) is proposed wherein an attempt is made to speed up convergence by increasing or decreasing the mutability of offspring proportional to the increase in fitness from the last iteration (Zheng et al., 2014). If the current best individual increases in fitness during an iteration, the mutability of the next generation is increased. Conversely, when the fitness does not increase, the mutability of new individuals is decreased. Offspring

generated by other parents are generated according to earlier EFWA-methods. Furthermore, the study argues that Gaussian sparks can be removed from the algorithm entirely, as in most results they will not increase diversity. Again, this results in equal or sometimes better performance than EFWA while reducing runtime.

Likewise, a variant of the original algorithm has been designed for PPA; a variant named MPPA (Modified PPA), as presented in Sulaiman et al. (2014a). MPPA uses a fixed number of offspring instead of using the relative fitness of an individual to calculate its number of offspring. Furthermore, new individuals are generated one by one, where if it is not better than its parent, it is discarded and a new individual is generated. This is repeated up to three times, each time using a different formula when generating the position of the new individual. MPPA shows better performance than PPA and the Artificial Bee Colony Algorithm (Sulaiman et al., 2014a).

Although there are many published adaptations on the original algorithms, the comparison in this study is made between the original descriptions of FWA and PPA only. Both algorithms are implemented from scratch and have their performance compared on a specific set of benchmark test functions. But first, the algorithms will be functionally aligned for qualitative comparison.

2 FIREWORKS & PLANT PROPAGATION

Although FWA and PPA abide by very similar algorithmic philosophies, equations and routines are deployed differently. A perfect example of this can be found in the order of procedures which are common for all population-based algorithms: initialization, offspring generation, and selection of individuals. FWA starts with initialization, then loops over generating offspring and selecting individuals respectively. PPA also starts with initialization, but then loops over selecting individuals and generating offspring. Although the order of these procedures is different for FWA and PPA, they can be interchanged without loss of generality. For a uniformly formatted overview of both algorithms, the reader is referred to table 1. In order to facilitate comparison and ease of reading the formulas and terminology are linguistically generalized, but otherwise stay functionally consistent with the seminal papers (Tan and Zhu, 2010; Salhi and Fraga, 2011).

Initialization of the population, as done in the seminal papers (Tan and Zhu, 2010; Salhi and Fraga, 2011), is nearly identical for FWA and PPA. Individuals in the initial population are generated by drawing

Table 1: Algorithmic subroutines of FWA and PPA where N represents the size of the current population, x_{max} and x_{min} the current best and worst individual, $r \in [0, 1]$ a uniform random number, x_{ij} the position of individual i in dimension j , the benchmark function f , the number of dimensions in the benchmark function D , and a_j and b_j the lower and upper bound of the benchmark function in dimension j . Algorithm specific parameters are: $PopSize$ the number of individuals in a population after selection, \hat{m} the parameter that controls the number of offspring per individual, m the parameter that controls the number of Gaussian Sparks, n_{max} and n_{min} the maximum and minimum number of offspring per individual, and d_{max} the maximum distance of offspring.

Subroutine	FWA	PPA
Initialization of the population	$PopSize$ individuals uniform random over whole domain	$PopSize$ individuals uniform random over whole domain
Assigning fitness	$F_1(x_i) = \frac{f(x_{max}) - f(x_i) + \epsilon}{\sum_{i=1}^N (f(x_{max}) - f(x_i)) + \epsilon}$ $F_2(x_i) = \frac{f(x_i) - f(x_{min}) + \epsilon}{\sum_{i=1}^N (f(x_i) - f(x_{min})) + \epsilon}$	$z(x_i) = \frac{f(x_{max}) - f(x_i)}{f(x_{max}) - f(x_{min})}$ $F(x_i) = \frac{1}{2}(\tanh(4 \cdot z(x_i) - 2) + 1)$
Number of offspring	$\hat{n}(x_i) = \hat{m}F_1(x_i)$ $n(x_i) = \begin{cases} n_{min} & \text{if } \hat{n}(x_i) < n_{min} \\ n_{max} & \text{if } \hat{n}(x_i) > n_{max} \\ round(\hat{n}(x_i)) & \text{otherwise} \end{cases}$	$n(x_i) = \lceil n_{max}F(x_i)r \rceil$
Position of offspring	$d(x_i) = 2(r - 0.5)F_2(x_i)d_{max}$ for $\lceil D \cdot r \rceil$ random dimensions, do: $x_{ij}^* = x_{ij} + d(x_i)$	$d_j(x_i) = 2(r - 0.5)(1 - F(x_i))$ $x_{ij}^* = x_{ij} + (b_j - a_j)d_j(x_i)$
Generating additional (Gaussian) offspring	m extra individuals	-
Position of additional (Gaussian) offspring	$g = \mathcal{N}(\mu, \sigma^2), \text{ with } \mu = \sigma = 1$ for $\lceil D \cdot r \rceil$ random dimensions, do: $x_{ij}^* = x_{ij} * g$	-
Bounds correction	$x_{ij}^* = \begin{cases} x_{ij}^* & \text{if } a_j \leq x_{ij}^* \leq b_j \\ a_j + x_{ij}^* \bmod (b_j - a_j) & \text{otherwise} \end{cases}$	$x_{ij}^* = \begin{cases} a_j & \text{if } x_{ij}^* < a_j \\ b_j & \text{if } x_{ij}^* > b_j \\ x_{ij}^* & \text{otherwise} \end{cases}$
Selection of individuals	Best individual and $PopSize - 1$ additional individuals which are selected through: $R(x_i) = \sum_{j \in N} x_i - x_j $ $p_{selection}(x_i) = \frac{R(x_i)}{\sum_{j \in N} R(x_j)}$	Best $PopSize$ individuals are selected

a uniform random number from the entire domain for each of the dimensions. Although in the seminal paper on FWA initialization is done on either $[30, 50]^D$ or $[15, 30]^D$, personal correspondence with the author, Ying Tan, confirmed that this could just as well be done on the entire domain, ultimately making it identical to PPA's initialization method.

Assigning fitness for both algorithms is done by 'normalizing' the fitness values for all individuals in the population to lie in $(0, 1)$. One big difference between FWA and PPA is that FWA defines *two different* formulas for the calculation of relative fitness for each individual: $F_1(x_i)$ and $F_2(x_i)$. The outcome of F_1 increases proportionately from the best to the worst individual and is used when calculating the number of offspring that an individual generates. Contrarily, the outcome of F_2 decreases proportionate from the best to the worst individual and is used when deciding the mutation size of the newly generated offspring. To-

gether, the two formulas form the essence of FWA's paradigm; good individuals produce many offspring with small mutations, and bad individuals vice versa.

Although it abides by the same paradigm, PPA defines the normalized or 'relative' fitness values with just one formula: $F(x_i)$. The outcome of this formula is then used to calculate the number of offspring for an individual, where higher values result in a higher number of offspring. Conversely, the inverse $(1 - F(x_i))$ is proportionate to the mutation size of the offspring. However, unlike FWA, the relative fitness is calculated in two steps: the individual's objective values are first normalized to be in $[0, 1]$, after which a 'mapping function' (a hyperbolic tangent) ensures that values lie strictly in $(0, 1)$. As the authors note: this is necessary, because the methods that are used to calculate the offspring distance would otherwise produce zero, resulting in offspring that is at the exact same location as the parent. According to the authors, this

mapping function also “[further emphasizes better solutions over those which are not as good]” (Salhi and Fraga, 2011). Also, note that FWA’s relative fitnesses are proportionate whereas PPA’s relative fitness is linear before applying the hyperbolic tangent; FWA’s number of offspring and mutation are only somewhat inversely related, whereas PPA’s number of offspring and mutation distance are directly inversely related.

Generating offspring is done differently in both algorithms. FWA uses two methods, the first (‘regular’) method produces offspring for each individual, in numbers proportionate to that individual’s F_1 -fitness, and their mutability proportionate to the somewhat inversely related F_2 -fitness. FWA’s additional method of offspring generation (“Gaussian sparks”) selects m individuals randomly to generate one *extra* offspring. According to the authors, “this improves diversity among individuals” (Tan and Zhu, 2010).

Contrarily, PPA only has one method for generating offspring which is relatively similar to the first method of FWA. Each individual produces offspring, the number of which is proportionate to its fitness, and the mutability of which is proportionate to the *direct inverse* of its fitness.

In FWA, the **number of offspring** is determined for each individual by multiplying its F_1 -fitness with \hat{m} , a variable that determines the maximum number of offspring and rounds the outcome. Then, this outcome is verified to be between n_{max} and n_{min} , parameters that limit the number of offspring per individual. If a parameter is exceeded, the value of the number of offspring is truncated to the respective parameter. For the “Gaussian sparks”-method, FWA randomly selects m individuals from the population that each generate *one* extra offspring. Note that FWA uses four parameters to ensure that the number of offspring is within desirable range: a parameter \hat{m} that controls the number of offspring per individual, number of Gaussian sparks m , minimum and maximum number of offspring n_{min} and n_{max} . Remarkably, these four parameters could be reduced to three by only defining the number of Gaussian sparks m , the minimum number of offspring n_{min} , and the maximum number of offspring n_{max} and then replacing m in the equation by a factor of n_{max} .

In PPA, the number of offspring for an individual is determined by multiplying its fitness $F(x_i)$ with a random number $r \in [0, 1)$ and the maximum number of offspring per individual n_{max} , and then rounding up the outcome. Note that PPA therefore requires only one parameter to keep the number of offspring within a desirable range.

Calculating the **mutation of offspring** is another aspect in which both algorithms are just slightly different. FWA uses the F_2 -fitness and multiplies this

with a maximum mutability (d_{max}). The resulting mutation is then applied to a number of dimensions z , where z is found by multiplying a uniform random number $r \in [0, 1)$ by the number of dimensions D and then rounding up. Additionally, FWA randomly selects m ‘Gaussian candidates’ that will generate one extra offspring. This offspring is mutated by multiplying z dimensions with a Gaussian factor $g = \mathcal{N}(1, 1)$. Again, according to the authors, this method “ensures diversity among individuals” (Tan and Zhu, 2010).

PPA first calculates a mutation by taking the inverse relative fitness $1 - F(x_i)$, which is then multiplied with a random number in $[-0.5, 0.5]$ and scaled to the size of the domain of the objective function. Unlike the predetermined maximum amplitude (d_{max}) found in FWA, this construct ensures that no assumptions are made about the domain of the objective function. Also note that whereas FWA reuses a single mutation size ($d(x_i)$) for each of the z selected dimensions, PPA calculates a separate mutation size ($d_j(x_i)$) with a new random number for each dimension. Thus, the difference between the algorithms is that FWA mutates *some* dimensions *equally* whereas PPA mutates *all* dimensions *differently*. Possibly, the best mutation operator would be to mutate some dimensions differently – a combination of the two methods described.

Bounds correction is required for both algorithms to ensure that generated offspring do not generate outside the bounds of the benchmark function. Neither algorithm prevents this from happening, but instead check whether newly generated individuals are outside of the specified boundaries. If an individual would be generated outside of these boundaries, the location of the individual is adjusted to be within the boundaries. FWA does this by re-mapping the value to be within the bounds with a correction function (see table 1). Individuals outside bounds will be adjusted to a value that is no bigger than the difference between the lower and upper bound, after which the value is added to the lower bound. Most notably, this function tends to correct the values close to the central point between the domain’s bounds.

PPA corrects individuals by simply truncating any out-of-bound value to the respective bound that was exceeded. This is easier to compute than FWA’s method, but the authors correctly identify a possible bias: “there will be some preference for points being generated at the bounds of the search space” Salhi and Fraga (2011). In short, both algorithms appear to be biased in keeping their offspring with the problem domain’s bounds.

Although **selection of individuals** for FWA and PPA is done differently, both algorithms utilize an elitist approach; the best individual in the population is

Table 2: List of the benchmark test functions used in this study.

Benchmark name	Function	Bounds	Global Minimum
[2D] Six-Hump-Camel	$f(x_1, x_2) = (4 - 2.1x_1^2 + \frac{x_1^4}{3})x_1^2 + x_1x_2 + (-4 + 4x_2^2)x_2^2$	$x_1 \in [-3, 3],$ $x_2 \in [-2, 2]$	$f(x^*) = -1.032,$ at $x^* = (0.0898, -0.7126)$ and $(-0.0898, 0.7126)$
[2D] Martin-Gaddy	$f(x_1, x_2) = (x_1 - x_2)^2 + (\frac{x_1 + x_2 - 10}{3})^2$	$x_1, x_2 \in [-20, 20]$	$f(x^*) = 0, \text{ at } x^* = (5, 5)$
[2D] Goldstein-Price	$f(x_1, x_2) = (1 + (x_1 + x_2 + 1)^2 (19 - 14x_1 + 3x_1^2 - 14x_2 + 6x_1x_2 + 3x_2^2)) (30 + (2x_1 - 3x_2)^2 (18 - 32x_1 + 12x_1^2 + 48x_2 - 36x_1x_2 + 27x_2^2))$	$x_1, x_2 \in [-2, 2]$	$f(x^*) = 3, \text{ at } x^* = (0, -1)$
[2D] Branin	$f(x_1, x_2) = (x_2 - \frac{5.1}{4\pi^2}x_1^2 + \frac{5}{\pi}x_1 - 6)^2 + 10(1 - \frac{1}{8\pi})\cos(x_1) + 10$	$x_1, x_2 \in [-5, 15]$	$f(x^*) = 0.3979,$ at $x^* = (-\pi, 12.275),$ $(\pi, 2.275)$ and $(9.42478, 2.475)$
[2D] Easom	$f(x_1, x_2) = -\cos(x_1)\cos(x_2)e^{(-(x_1-\pi)^2-(x_2-\pi)^2)}$	$x_1, x_2 \in [-100, 100]$	$f(x^*) = -1, \text{ at } x^* = (\pi, \pi)$
Rosenbrock	$f(\mathbf{x}) = \sum_{i=1}^{d-1} (100(x_{i+1} - x_i^2)^2 + (x_i - 1)^2)$	$x \in [-5, 10]^D$	$f(x^*) = 0, \text{ at } x^* = (1)^D$
Ackley	$f(\mathbf{x}) = -20 \cdot \exp(-0.2 \cdot \sqrt{\frac{1}{d} \sum_{i=1}^d x_i^2}) - \exp(\frac{1}{d} \sum_{i=1}^d \cos(2\pi x_i)) + 20 + e$	$x \in [-100, 100]^D$	$f(x^*) = 0, \text{ at } x^* = (0)^D$
Griewank	$f(\mathbf{x}) = 1 + \sum_{i=1}^d \frac{x_i^2}{4000} - \prod_{i=1}^d \cos(\frac{x_i}{\sqrt{i}})$	$x \in [-600, 600]^D$	$f(x^*) = 0, \text{ at } x^* = (0)^D$
Rastrigrin	$f(\mathbf{x}) = 10d + \sum_{i=1}^d (x_i^2 - 10\cos(2\pi x_i))$	$x \in [-5.12, 5.12]^D$	$f(x^*) = 0, \text{ at } x^* = (0)^D$
Schwefel	$f(\mathbf{x}) = 418.9829d - \sum_{i=1}^d (x_i \sin(\sqrt{ x_i }))$	$x \in [-500, 500]^D$	$f(x^*) = 0, \text{ at } x^* = (420.9687)^D$
Ellipse	$f(\mathbf{x}) = \sum_{i=1}^d 10000(\frac{i-1}{d-1})x_i^2$	$x \in [-100, 100]^D$	$f(x^*) = 0, \text{ at } x^* = (0)^D$
Cigar	$f(\mathbf{x}) = x_1^2 + \sum_{i=2}^d 10000x_i^2$	$x \in [-100, 100]^D$	$f(x^*) = 0, \text{ at } x^* = (0)^D$
Tablet	$f(\mathbf{x}) = 10000x_1^2 + \sum_{i=2}^d x_i^2$	$x \in [-100, 100]^D$	$f(x^*) = 0, \text{ at } x^* = (0)^D$
Sphere	$f(\mathbf{x}) = \sum_{i=1}^d x_i^2$	$x \in [-100, 100]^D$	$f(x^*) = 0, \text{ at } x^* = (0)^D$

always transferred to the new population. FWA uses a selection procedure wherein the current best individual is selected first, and subsequent individuals are chosen through ‘dispersity proportionate selection’. First, each individual’s dispersity is calculated by the total Euclidean distance the individual is away from all other individuals. Then, a dispersity-proportional probability is assigned to each individual, and *PopSize* individuals are randomly selected for the new population. According to the authors, this process “ensures diversity among individuals” (Tan and Zhu, 2010). Contrarily, PPA simply selects the *PopSize* best individuals for the new population.

A final implementational difference between FWA and PPA is the method the algorithms use to prevent division by zero errors. Division by zero can, in the case of both of these algorithms, be caused by having a population where all individuals have the exact same fitness value. To address this problem, FWA uses the machine epsilon (the smallest representable number ϵ in the computer such that $1 + \epsilon > 1$) in the numerator and the denominator, and as a result, gives each indi-

vidual in a homogeneous population a relative fitness of 1. Whereas PPA uses an if-statement wherein it checks if the maximum is equal to the minimum. If this is the case, it will assign a relative fitness of 0.5 to each individual in the population.

3 METHOD

Benchmark test functions behave notoriously fickle across the body of literature. Functional descriptions, domain ranges, vertical scaling, initialization values and even for the exact spelling of a function’s name, a multitude of alternatives can be found. It is therefore wise to use explicit definitions. The benchmark test functions used in this study are a union of the benchmark test functions listed in the seminal papers (Tan and Zhu, 2010; Salhi and Fraga, 2011) and are listed in table 2. A selection of graphs of the 2-dimensional variants of these functions can be found in 1.

Comparison of benchmark scores is done through observing the number of evaluations, not the number of

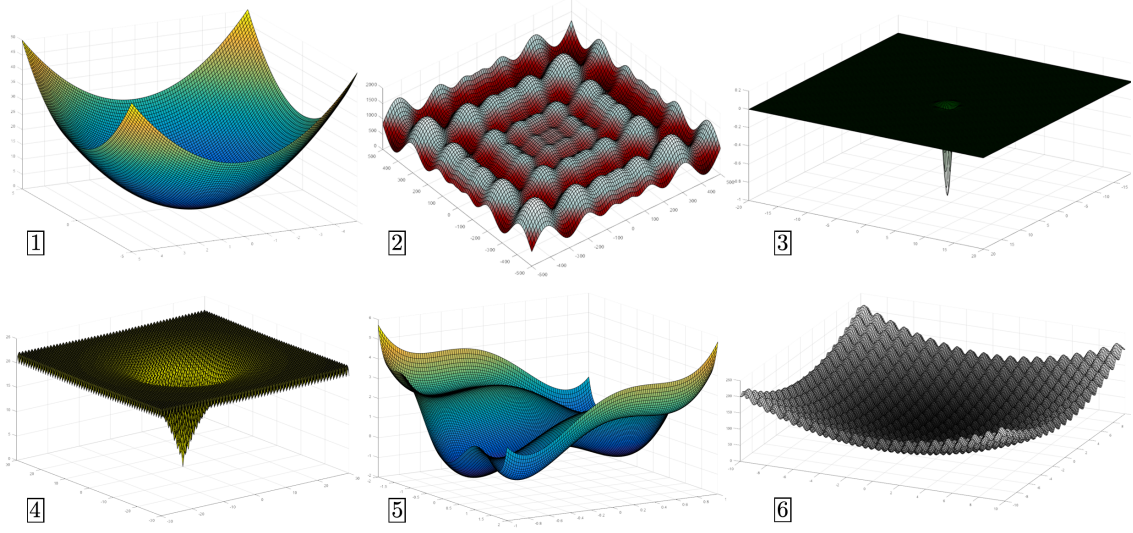


Figure 1: Six of the fourteen benchmark test functions test used in this study: (1) Sphere, (2) Schwefel, (3) Easom, (4) Ackley, (5) Six-Hump-Camel, and (6) Rastrigrin. Of these, only Easom and Six-Hump are 2D only; the others are $D \in [2, 100]$. All functions are associated with minimization tasks; the best objective value (vertical axes) is located lowest; exact formulae and locations and values for global minima can be found in table 2.

generations or total processing time. Both algorithms are run for 10,000 function evaluations for twenty trials on each of the benchmark test functions for all available dimensions (limited to 100 dimensions for the N-dimensional benchmark test functions). The best objective value from the population for each run at each function evaluation is recorded. Attained values are normalized by adding a constant to the benchmark function such that its global minimum has an exact value of zero, thereby enabling a logarithmic vertical axis in the result graphs and facilitating direct comparison between different benchmark functions.

All experiments are performed using the parameters from the original papers of both algorithms. For FWA, $PopSize = 30$, the number of offspring per individual $\hat{m} = 50$, the number of Gaussian sparks $m = 5$, the minimum and maximum number of offspring per individual $n_{min} = 2$ and $n_{max} = 40$, and the maximum mutability $d_{max} = 40$. For PPA these are $PopSize = 30$, and the maximum number of offspring $n_{max} = 5$.

The study by Zheng et al. (2013) indicates that the quality of results of FWA deteriorates when applied to benchmark test functions that are shifted away from the origin. This can be tested by translating both the function and its domain for the experiments (such as they are described above). By translating both the function and the domain itself, none of the output values of the benchmark test function change position relative to its bounds. For example, with a translation of -10 , the sphere function would become $f(\mathbf{x}) = \sum_{i=1}^d (x_i + 10)^2$

with bounds $x \in [-110, 90]^D$ by which the global optimum relocates to $f(x^*) = 0$ at $x^* = (-10)^D$ without changing its relative position in the domain or changing the range of the function. Of the five 2D functions in the set, none have their global minimum at the origin. For these, different shift values are applied to shift the global minimum *onto* the origin, resulting in ‘centered’ functions. A set of six translations of different magnitude is applied to all centered 2D functions. Experiments are repeated for all N-dimensional benchmark test functions for all available dimensions, where each of the dimensions is translated with the found translation value.

Implementation of all benchmark test functions, graphs, FWA, and PPA is done using Python 3.6. Code for replicating this study is available at <https://github.com/WouterVrielink/FWAPPA>.

4 RESULTS

On eight out of nine non-shifted N-dimensional benchmark test functions, FWA generally attains the best normalized objective values (figure 2). Only for the Schwefel function – and Rosenbrock and Griewank in the lower dimensions – PPA shows better results. Note that FWA only shows low variance in the best found objective values for the Rosenbrock and Schwefel function. On other functions, the best found objec-

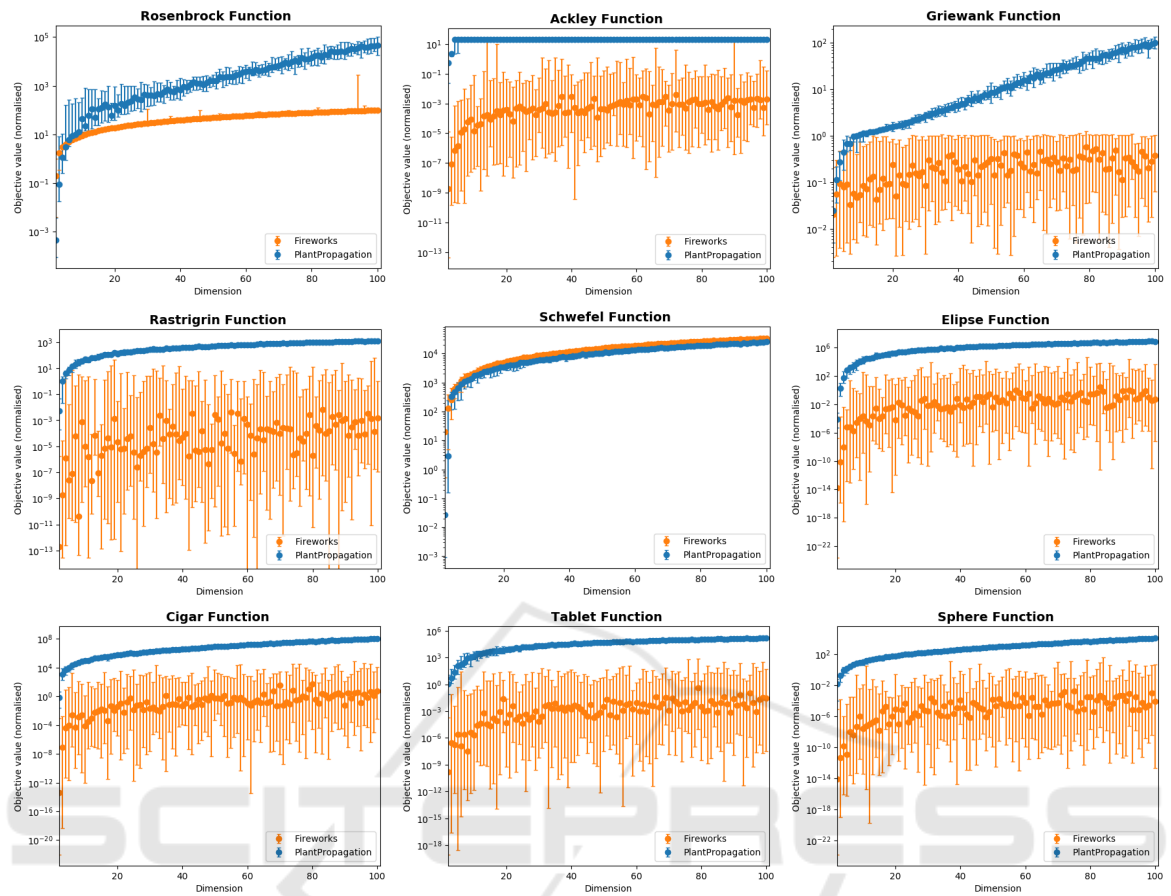


Figure 2: Normalized benchmark scores for the nine N-dimensional benchmark test functions. The horizontal axes show the number of dimensions of the benchmark function, while the vertical axes show the median best values found after 10,000 evaluations ($N = 20$). The error bars show the lowest and highest achieved results (0th and 100th percentile).

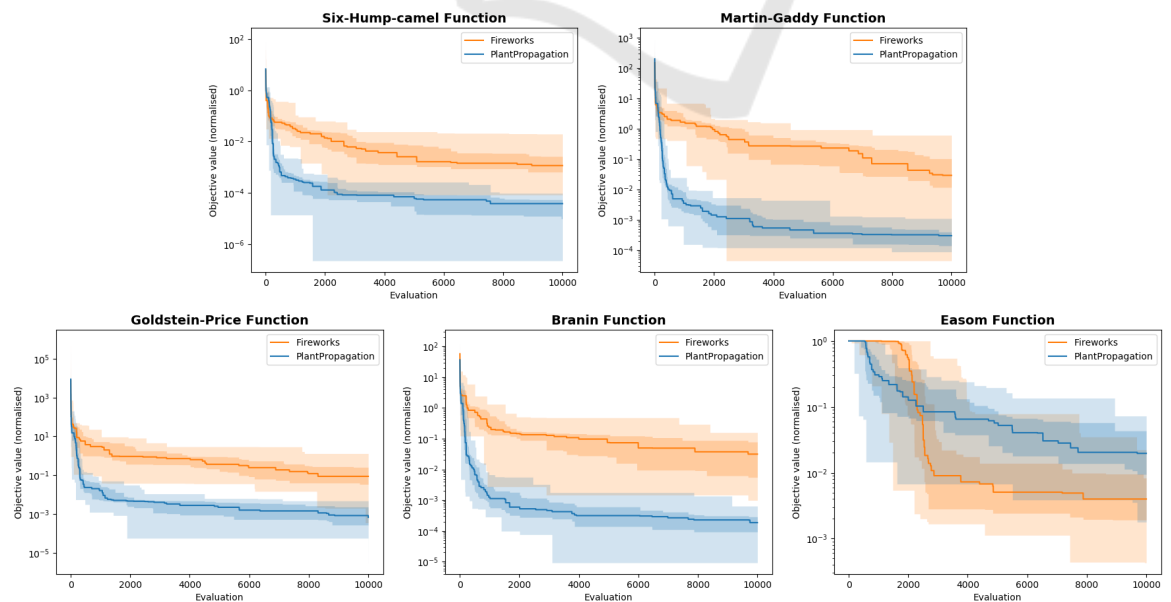


Figure 3: Results for the five 2-dimensional benchmark test functions. The horizontal axes show the number of evaluations of the benchmark function, while the vertical axes show the median best values found ($N = 20$). Filled areas show quartiles.

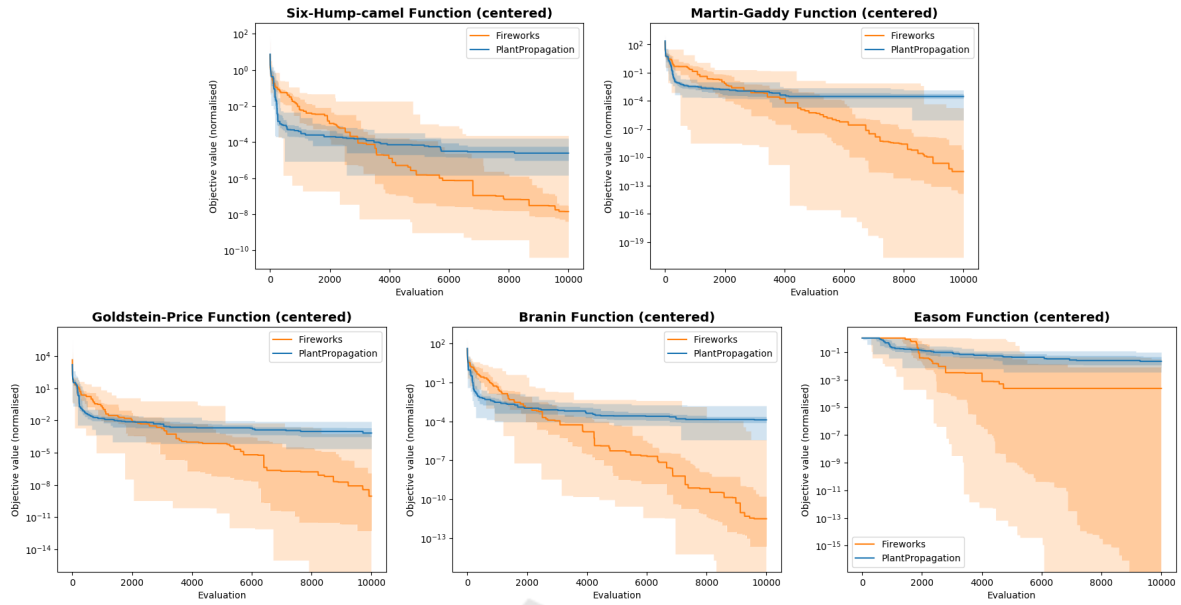


Figure 4: Results for the five centered 2-dimensional benchmark test functions. The horizontal axes show the number of evaluations of the benchmark function, while the vertical axes show the median best values found ($N = 20$). Lower values are considered better. The filled areas show the quartiles.

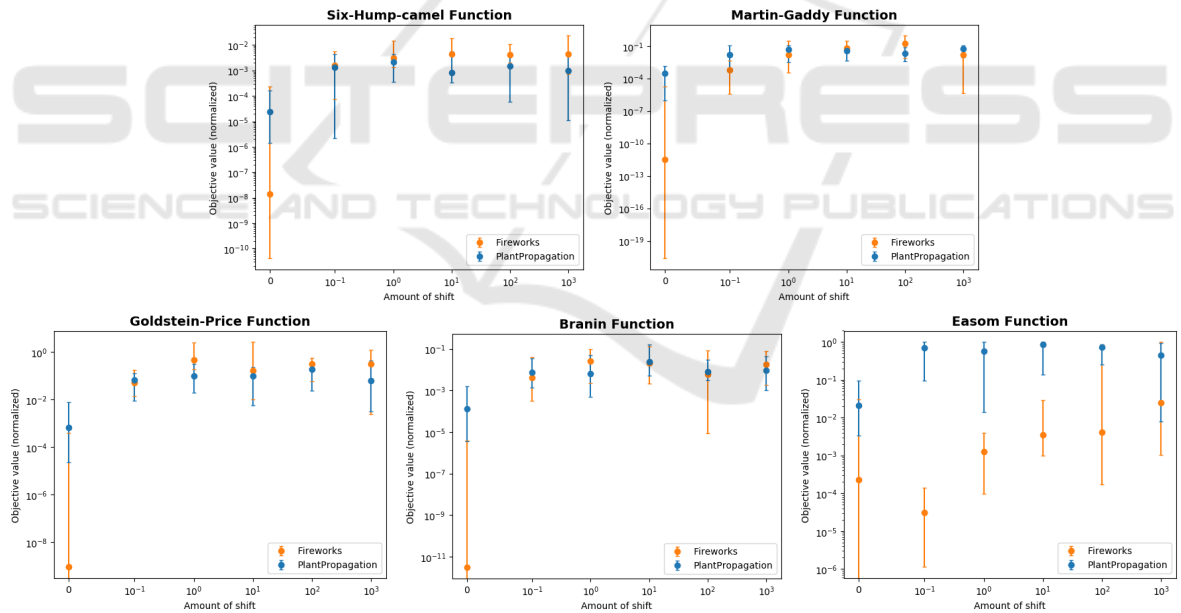


Figure 5: Results for the five 2-dimensional benchmarks when they are translated with different values. The horizontal axes show the amount of translation of the benchmark function, while the vertical show the median best values found ($N = 20$). Lower values are considered better. The error bars show the lowest and highest achieved results (0th and 100th percentile).

tive values of FWA have high variance and results can vary as much as a factor 10^{10} between trials. Results of PPA generally have lower variance, but it seems to be affected more by an increase in dimensionality than FWA. PPA is unable to find good values on the Ackley function for everything but the lowest dimensions.

Figure 3 shows normalized benchmark scores on the five 2D functions, where PPA outperforms FWA in four out of five cases. Both algorithms generally converge very quickly in the first 1,000 function evaluations, after which convergence speed slows down considerably. In figure 4 the results for the same functions with the minimum shifted onto the origin ('centered')

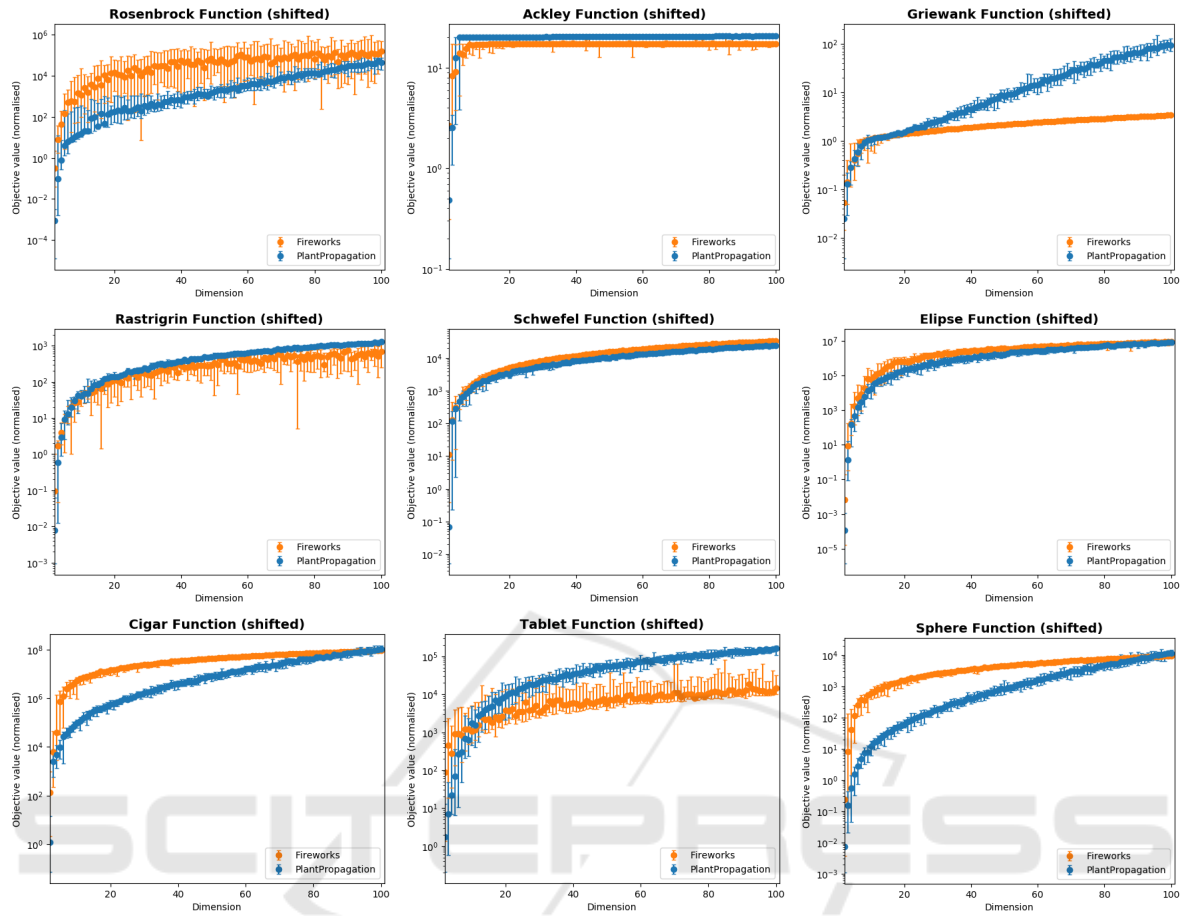


Figure 6: Normalized benchmark scores for the nine N-dimensional benchmark test functions when shifted such that the global minimum is at $x^* = (-10)^D$. The horizontal axes show the number of dimensions of the benchmark function, while the vertical axes show the median best values found after 10,000 evaluations ($N = 20$). Lower values are considered better. The error bars show the lowest and highest achieved results (0th and 100th percentile).

are shown. In this second figure, FWA outperforms PPA, but the high variance of FWA reappears. Statistical analysis using the Mann-Whitney U test confirms that the newly obtained objective values are significantly different for FWA, whereas those for PPA there is no significant difference (FWA gives $p < 0.05$, while PPA gives $p > 0.05$).

To see the effect of the translation distance, a set of six translations of different magnitudes is applied to all centered 2-dimensional functions. Results are shown in figure 5. As the function's global minimum benchmark is shifted further away from the origin, the performance of FWA deteriorates, while PPA does not seem to be affected. The effects of the translation appear to be the least significant on Easom, where the performance of FWA starts deteriorating with translations greater than 1. Note that, like shown in figures 2, 3, and 4, the variance of FWA decreases when the global optimum is shifted away from the origin. Otherwise, the variance of FWA is similar to PPA's. The

astute observer might notice that PPA's performance also deteriorates slightly when the global minimum of the benchmark function is shifted away from the origin – which would indicate a locational bias – but this difference cannot be proven to be significant.

For the translated N-dimensional benchmark test functions, a translation of +10 was used over all dimensions (figure 6). Compared to the results in figure 2, results of both algorithms are much more similar and FWA achieves better results than PPA on four out of nine benchmark test functions. Most noteworthy is the visual difference in the results of FWA between the untranslated and the translated benchmark test functions in figures 2 and 6 respectively. This difference is not apparent in the results of PPA. This is confirmed for all benchmark functions except Rosenbrock and Schwefel through statistical analysis using the Mann-Whitney U test. On Rosenbrock, Rastrigrin and Tablet, FWA again seems to have more variability in its end results. On the other functions however, the variance

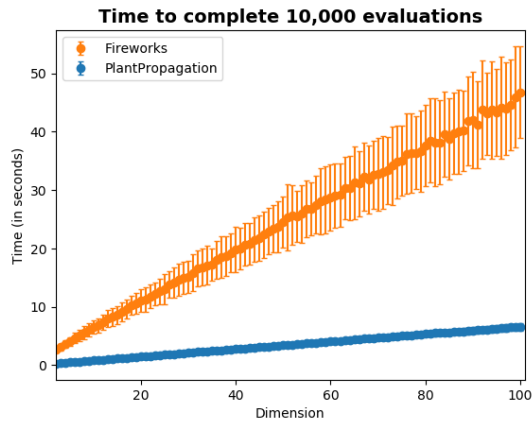


Figure 7: Average computing time in seconds to complete 10,000 evaluations in FWA and PPA ($N = 20$). Average is taken only from the N -dimensional benchmarks without applying translation. The horizontal axis shows the number of dimensions of the benchmark test functions, and the vertical axis the average time taken. The error bars show one standard deviation.

of FWA now seems to be similar and sometimes even smaller than that of PPA. Furthermore, FWA seems to generally scale better than PPA when the dimensionality of the problem increases.

A final notable difference is the computing times of the algorithms (figure 7). The runtime for 10,000 evaluations averaged over all N -dimensional benchmark test functions ($N = 20$) of FWA and PPA has been compared by fitting a straight line. Fits were tight with a correlation coefficient of 0.927 and a standard error of 1.371×10^{-3} for FWA and a correlation coefficient of 0.980 and a standard error of 9.998×10^{-5} for PPA. Thus, both algorithms increase linearly with the dimensionality of the benchmark test function. However, gradients were 0.441 for FWA and 0.064 for PPA, indicating that FWA slows down almost seven times as fast when the number of dimensions is increased. The intercepts were 2.104 and 0.173 for FWA and PPA respectively, indicating that FWA takes significantly more time to create and evaluate an initial population.

5 DISCUSSION

Although FWA and PPA have quite similar algorithmic paradigms, their performance is quite different. FWA outperforms PPA as long as a function's minimum is at the origin, but things notably change when this is not the case. In general, it looks like PPA is more suitable for low-dimensional problems, but again, the location of the global minimum seems to be of crucial importance. An avenue of further research might involve further exploring the influence of the *magnitude*

of these shifts. However, the magnitude of the distance between the global minimum and the origin seems to have little effect. Generally speaking, it seems like FWA is more suitable for situations where dimensionality is high, whereas for low-dimensional problems PPA is the better option.

Especially for FWA, multiple repetitions are advised when searching for the global optimum, as results might have high variance. It appears to converge faster than PPA in most cases, but this is at the cost of an increase in computational costs. Possible causes of the high computational cost of FWA are the sub-routines for generating Gaussian offspring, bounds correction (using modulo), and the selection of individuals ($O(n^2)$, versus $O(1)$ in PPA). In some real-life problems the computational cost of evaluating the objective function is multiple orders of magnitude larger than the computational cost of the algorithm (Keane, 1996; Ygge and Akkermans, 1996). In this case, the most important property of a combinatorial optimization algorithm is that it converges in as few evaluations as possible.

Arguably, the most interesting result of this study is the preference of FWA for centered benchmark test functions. There seems to be a universal pattern: when a function has its global minimum (further) away from the origin, PPA performs better and FWA behaves less erratic. The performance of FWA deteriorates significantly within translations within the size of 1, and larger translations seem to not have much more effect. The preference of FWA for benchmark test functions that have a global minimum positioned on the origin was also observed by Zheng et al. (2013), but while improvements were made in that study, the exact behaviour of this bias has not yet been studied. To this end, a more general method for detecting biases in continuous optimization algorithm needs to be developed, or as Sörensen (2015) states: "Perhaps a set of tools is needed, i.e., a collection of statistical programs or libraries specifically designed to determine the relative quality of a set of algorithms on a set of problem instances."

The source code for algorithms, benchmark test functions, statistical methods, and graphs used in this study is provided on GitHub (Vrieling, 2019) and may be freely used by anyone who wishes to compare the results of these or other optimization algorithms.

ACKNOWLEDGEMENTS

The authors would like to thank Marcus Pfundstein (former student, UvA) for his preliminary work compiling a set of benchmarks, Quinten van der Post (col-

league, UvA) for allowing us to use his server as a computing platform, Hans-Paul Schwefel (Professor Emeritus, University of Dortmund) for answering our questions about the Schwefel benchmark function by email, and Ying Tan (Professor, Peking University), Abdellah Salhi (Professor, University of Essex), and Eric Fraga (Professor, UCL London) for answering our questions on FWA and PPA respectively.

REFERENCES

- Alba, E. and Dorronsoro, B. (2005). The exploration/exploitation tradeoff in dynamic cellular genetic algorithms. *IEEE transactions on evolutionary computation*, 9(2):126–142.
- Audibert, J.-Y., Munos, R., and Szepesvári, C. (2009). Exploration–exploitation tradeoff using variance estimates in multi-armed bandits. *Theoretical Computer Science*, 410(19):1876–1902.
- Cheraitia, M., Haddadi, S., and Salhi, A. (2017). Hybridizing plant propagation and local search for uncapacitated exam scheduling problems. *International Journal of of Services and Operations Management*.
- El-Beltagy, M. A. and Keane, A. J. (2001). Evolutionary optimization for computationally expensive problems using gaussian processes. In *Proc. Int. Conf. on Artificial Intelligence*, volume 1, pages 708–714. CiteSeer.
- Geem, Z. W., Kim, J. H., and Loganathan, G. V. (2001). A new heuristic optimization algorithm: harmony search. *simulation*, 76(2):60–68.
- Geleijn, R., van der Meer, M., van der Post, Q., and van den Berg, D. (2019). The plant propagation algorithm on timetables: First results. In *EvoStar 2019 “The Leading European Event on Bio-Inspired Computation”*.
- Glover, F. (1989). Tabu search—part i. *ORSA Journal on computing*, 1(3):190–206.
- Glover, F. (1990). Tabu search—part ii. *ORSA Journal on computing*, 2(1):4–32.
- Imran, A. M. and Kowsalya, M. (2014). A new power system reconfiguration scheme for power loss minimization and voltage profile enhancement using fireworks algorithm. *International Journal of Electrical Power & Energy Systems*, 62:312–322.
- Ishii, S., Yoshida, W., and Yoshimoto, J. (2002). Control of exploitation–exploration meta-parameter in reinforcement learning. *Neural networks*, 15(4-6):665–687.
- Keane, A. (1996). The design of a satellite beam with enhanced vibration performance using genetic algorithm techniques. *Journal of the Acoustical Society of America*, 99(4):2599–2603.
- Kirkpatrick, S., Gelatt, C. D., and Vecchi, M. P. (1983). Optimization by simulated annealing. *science*, 220(4598):671–680.
- Li, J., Zheng, S., and Tan, Y. (2014). Adaptive fireworks algorithm. In *Evolutionary Computation (CEC), 2014 IEEE Congress on*, pages 3214–3221. IEEE.
- Moshrefi-Torbati, M., Keane, A., Elliott, S., Brennan, M., and Rogers, E. (2003). Passive vibration control of a satellite boom structure by geometric optimization using genetic algorithm. *Journal of Sound and Vibration*, 267(4):879–892.
- Paauw, M. and van den Berg, D. (2019). Paintings, polygons and plant propagation. In *International Conference on Computational Intelligence in Music, Sound, Art and Design (Part of EvoStar)*, pages 84–97. Springer.
- Salhi, A. and Fraga, E. S. (2011). Nature-inspired optimisation approaches and the new plant propagation algorithm.
- Selamoğlu, B. İ. and Salhi, A. (2016). The plant propagation algorithm for discrete optimisation: The case of the travelling salesman problem. In *Nature-inspired computation in engineering*, pages 43–61. Springer.
- Sörensen, K. (2015). Metaheuristics—the metaphor exposed. *International Transactions in Operational Research*, 22(1):3–18.
- Storn, R. and Price, K. (1997). Differential evolution—a simple and efficient heuristic for global optimization over continuous spaces. *Journal of global optimization*, 11(4):341–359.
- Sulaiman, M., Salhi, A., and Fraga, E. S. (2014a). The plant propagation algorithm: modifications and implementation. *arXiv preprint arXiv:1412.4290*.
- Sulaiman, M., Salhi, A., Selamoğlu, B. I., and Kirikchi, O. B. (2014b). A plant propagation algorithm for constrained engineering optimisation problems. *Mathematical Problems in Engineering*, 2014.
- Tan, Y. and Zhu, Y. (2010). Fireworks algorithm for optimization. In *International Conference in Swarm Intelligence*, pages 355–364. Springer.
- Vrieling, W. (2019). *FWA versus PPA*. <https://github.com/WouterVrieling/FWAPPA>.
- Weyland, D. (2015). A critical analysis of the harmony search algorithm-how not to solve sudoku. *Operations Research Perspectives*, 2:97–105.
- Ygge, F. and Akkermans, J. M. (1996). *Power load management as a computational market*. Höskolan i Karlskrona/Ronneby.
- Zheng, S., Janecek, A., Li, J., and Tan, Y. (2014). Dynamic search in fireworks algorithm. In *Evolutionary Computation (CEC), 2014 IEEE Congress on*, pages 3222–3229. IEEE.
- Zheng, S., Janecek, A., and Tan, Y. (2013). Enhanced fireworks algorithm.